

# Bios 6301: Assignment 3

Haoyang Yi

## Question 1

### 15 points

Write a simulation to calculate the power for the following study design. The study has two variables, treatment group and outcome. There are two treatment groups (0, 1) and they should be assigned randomly with equal probability. The outcome should be a random normal variable with a mean of 60 and standard deviation of 20. If a patient is in the treatment group, add 5 to the outcome. 5 is the true treatment effect. Create a linear model for the outcome by the treatment group, and extract the p-value (hint: see assignment1). Test if the p-value is less than or equal to the alpha level, which should be set to 0.05.

Repeat this procedure 1000 times. The power is calculated by finding the percentage of times the p-value is less than or equal to the alpha level. Use the `set.seed` command so that the professor can reproduce your results.

1. Find the power when the sample size is 100 patients. (10 points)

```
set.seed(414)
size = 100
times = 1000
res_pvalue = NULL
for(i in 1:times){
  group = sample(c(0,1),size,replace = T)
  y = rnorm(size,mean = 60,sd = 20)
  y[group==1]=y[group==1]+5
  mod = lm(y~group)
  res_pvalue[i] = coef(summary(mod))[2,4]
}
sum(res_pvalue < 0.05)/times # power when size = 100
```

```
## [1] 0.231
```

1. Find the power when the sample size is 1000 patients. (5 points)

```
set.seed(414)
size = 1000
times = 1000
res_pvalue = NULL
for(i in 1:times){
  group = sample(c(0,1),size,replace = T)
  y = rnorm(size,mean = 60,sd = 20)
  y[group==1]=y[group==1]+5
  mod = lm(y~group)
  res_pvalue[i] = coef(summary(mod))[2,4]
}
sum(res_pvalue < 0.05)/times # power when size = 1000
```

```
## [1] 0.974
```

## Question 2

### 14 points

Obtain a copy of the football-values lecture. Save the 2021/proj\_wr21.csv file in your working directory. Read in the data set and remove the first two columns.

1. Show the correlation matrix of this data set. (4 points)

```
library(MASS)
df = read.csv('proj_wr21.csv')
df = df[,-c(1,2)]
cor(df) # correlation matrix
```

##		rec_att	rec_yds	rec_tds	rush_att	rush_yds	rush_tds	fumbles
##	rec_att	1.0000000	0.9889836	0.9620513	0.2242480	0.2810831	0.2312038	0.6423627
##	rec_yds	0.9889836	1.0000000	0.9720400	0.2062038	0.2614786	0.2115013	0.6487247
##	rec_tds	0.9620513	0.9720400	1.0000000	0.2004448	0.2540571	0.2151580	0.6021914
##	rush_att	0.2242480	0.2062038	0.2004448	1.0000000	0.9779751	0.9308512	0.1446322
##	rush_yds	0.2810831	0.2614786	0.2540571	0.9779751	1.0000000	0.9298581	0.1761579
##	rush_tds	0.2312038	0.2115013	0.2151580	0.9308512	0.9298581	1.0000000	0.1809564
##	fumbles	0.6423627	0.6487247	0.6021914	0.1446322	0.1761579	0.1809564	1.0000000
##	fpts	0.9863078	0.9957911	0.9842850	0.2610623	0.3162080	0.2677821	0.6288445
##	fpts							
##	rec_att	0.9863078						
##	rec_yds	0.9957911						
##	rec_tds	0.9842850						
##	rush_att	0.2610623						
##	rush_yds	0.3162080						
##	rush_tds	0.2677821						
##	fumbles	0.6288445						
##	fpts	1.0000000						

1. Generate a data set with 30 rows that has a similar correlation structure. Repeat the procedure 1,000 times and return the mean correlation matrix. (10 points)

```
dfcor = cor(df)
dfcov = var(df)
dfmean = colMeans(df)
rescor = 0
set.seed(414)
# generate data from similar correlation structure
for (i in 1:1000){
  dt = mvrnorm(30,dfmean,dfcov)
  rescor = rescor+cor(dt)/1000
}
rescor
```

##		rec_att	rec_yds	rec_tds	rush_att	rush_yds	rush_tds	fumbles
##	rec_att	1.0000000	0.9884243	0.9609830	0.2219991	0.2781960	0.2294294	0.6353030
##	rec_yds	0.9884243	1.0000000	0.9711987	0.2055199	0.2600256	0.2106236	0.6424439
##	rec_tds	0.9609830	0.9711987	1.0000000	0.1997415	0.2527589	0.2145637	0.5957682
##	rush_att	0.2219991	0.2055199	0.1997415	1.0000000	0.9768849	0.9271844	0.1485957
##	rush_yds	0.2781960	0.2600256	0.2527589	0.9768849	1.0000000	0.9263911	0.1797179
##	rush_tds	0.2294294	0.2106236	0.2145637	0.9271844	0.9263911	1.0000000	0.1822552
##	fumbles	0.6353030	0.6424439	0.5957682	0.1485957	0.1797179	0.1822552	1.0000000
##	fpts	0.9856832	0.9956520	0.9838195	0.2591613	0.3136727	0.2658371	0.6226126
##	fpts							

```
## rec_att 0.9856832
## rec_yds 0.9956520
## rec_tds 0.9838195
## rush_att 0.2591613
## rush_yds 0.3136727
## rush_tds 0.2658371
## fumbles 0.6226126
## fpts 1.0000000
rescor=cor(df) # The difference between average cor matrix and true cor matrix is small.
```

```
##          rec_att    rec_yds    rec_tds    rush_att    rush_yds
## rec_att  6.661338e-16 -5.592727e-04 -1.068322e-03 -2.248962e-03 -2.887092e-03
## rec_yds  -5.592727e-04  6.661338e-16 -8.412918e-04 -6.839453e-04 -1.452956e-03
## rec_tds  -1.068322e-03 -8.412918e-04  6.661338e-16 -7.032827e-04 -1.298142e-03
## rush_att -2.248962e-03 -6.839453e-04 -7.032827e-04  6.661338e-16 -1.090128e-03
## rush_yds -2.887092e-03 -1.452956e-03 -1.298142e-03 -1.090128e-03  6.661338e-16
## rush_tds -1.774428e-03 -8.777031e-04 -5.942986e-04 -3.666769e-03 -3.467047e-03
## fumbles  -7.059659e-03 -6.280755e-03 -6.423238e-03  3.963492e-03  3.559944e-03
## fpts      -6.245745e-04 -1.391585e-04 -4.654498e-04 -1.900973e-03 -2.535274e-03
##          rush_tds    fumbles    fpts
## rec_att  -1.774428e-03 -7.059659e-03 -6.245745e-04
## rec_yds  -8.777031e-04 -6.280755e-03 -1.391585e-04
## rec_tds  -5.942986e-04 -6.423238e-03 -4.654498e-04
## rush_att -3.666769e-03  3.963492e-03 -1.900973e-03
## rush_yds -3.467047e-03  3.559944e-03 -2.535274e-03
## rush_tds  6.661338e-16  1.298764e-03 -1.945003e-03
## fumbles   1.298764e-03  6.661338e-16 -6.231894e-03
## fpts      -1.945003e-03 -6.231894e-03  6.661338e-16
```

### Question 3

21 points

Here's some code:

```
nDist <- function(n = 100) {
  df <- 10
  prob <- 1/3
  shape <- 1
  size <- 16
  list(
    beta = rbeta(n, shape1 = 5, shape2 = 45),
    binomial = rbinom(n, size, prob),
    chisquared = rchisq(n, df),
    exponential = rexp(n),
    f = rf(n, df1 = 11, df2 = 17),
    gamma = rgamma(n, shape),
    geometric = rgeom(n, prob),
    hypergeometric = rhyper(n, m = 50, n = 100, k = 8),
    lognormal = rlnorm(n),
    negbinomial = rnbinom(n, size, prob),
    normal = rnorm(n),
    poisson = rpois(n, lambda = 25),
    t = rt(n, df),
    uniform = runif(n),
```

```

    weibull = rweibull(n, shape)
  )
}
as.numeric(sapply(nDist(500), mean)[1])

```

```
## [1] 0.1004172
```

1. What does this do? (3 points)

```
round(sapply(nDist(500), mean), 2)
```

```
##          beta      binomial    chisquared    exponential          f
##          0.10          5.29          9.88          0.99          1.10
##          gamma    geometric hypergeometric    lognormal    negbinomial
##          0.99          1.84          2.56          1.69          31.37
##          normal      poisson          t          uniform      weibull
##          0.06          24.93          0.12          0.51          0.95
```

It return 2 digit means of each generated sample from beta,uniform... distributions, like mean of a generated sample with size 500 from beta distribution (5,45)

1. What about this? (3 points)

```
sort(apply(replicate(20, round(sapply(nDist(10000), mean), 2)), 1, sd))
```

```
##          beta      uniform          f    exponential          t
##    0.000000000    0.002236068    0.007677719    0.008750940    0.009947229
##          normal      weibull          gamma hypergeometric    binomial
##    0.010990426    0.012565617    0.012732057    0.013562720    0.015217718
##          geometric    lognormal    chisquared      poisson    negbinomial
##    0.021740092    0.024381831    0.046732723    0.047225662    0.079795792
```

It runs the simulation 20 times and calculates the standard deviation of 20 simulated means of each sample from beta,uniform..distributions .

In the output above, a small value would indicate that N=10,000 would provide a sufficient sample size as to estimate the mean of the distribution. Let's say that a value *less than 0.02* is "close enough".

2. For each distribution, estimate the sample size required to simulate the distribution's mean. (15 points)

```

b= nDist(50)
t = data.frame(distribution = names(b),true_mean = c(0.1,16/3,10,1,17/15,1,2,400/150,exp(1/2),32,0,25,0)
n = seq(10,2000,10) # starts with sample size = 10 and set gap = 10
set.seed(414)
diff = matrix(0,length(n),length(names(b)))
for (i in 1:length(names(b))) {
  for(j in 1:length(n))
    {diff[j,i]=as.numeric(sapply(nDist(n[j]), mean)[i])-t[i,2]}
}
diff = cbind(n,diff)
minval = NULL
for (i in 2:ncol(diff)){
  minval[i-1]=min(which(abs(diff[,i])<=0.02))
}
t$sizerequired = 10+10*(minval-1)
t

##          distribution true_mean sizerequired

```

```
## 1      beta  0.100000      10
## 2    binomial 5.333333     130
## 3   chisquared 10.000000    130
## 4   exponential 1.000000     90
## 5        f  1.133333     50
## 6      gamma 1.000000    200
## 7   geometric 2.000000    240
## 8 hypergeometric 2.666667    160
## 9    lognormal 1.648721     20
## 10  negbinomial 32.000000    100
## 11      normal 0.000000     80
## 12     poisson 25.000000    360
## 13        t  0.000000     70
## 14     uniform 0.500000     90
## 15     weibull 1.000000     10
```

#Don't worry about being exact. It should already be clear that  $N < 10,000$  for many of the distributions. You don't have to show your work. Put your answer to the right of the vertical bars (|) below.

distribution	N
beta	10
binomial	130
chisquared	130
exponential	90
f	50
gamma	200
geometric	240
hypergeometric	160
lognormal	20
negbinomial	100
normal	80
poisson	360
t	70
uniform	90
weibull	10