

homework7

Haoyang Yi

10/28/2022

Question 1

21 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {  
  if(exists(".Random.seed", envir = .GlobalEnv)) {  
    save.seed <- get(".Random.seed", envir = .GlobalEnv)  
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))  
  } else {  
    on.exit(rm(".Random.seed", envir = .GlobalEnv))  
  }  
  set.seed(n)  
  subj <- ceiling(n / 10)  
  id <- sample(subj, n, replace=TRUE)  
  times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))  
  dt <- as.POSIXct(sample(times, n), origin='2000-01-01')  
  mu <- runif(subj, 4, 10)  
  a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)  
  data.frame(id, dt, a1c)  
}  
x <- genData(500)
```

Perform the following manipulations: (3 points each)

1. Order the data set by id and dt.

```
x = arrange(x,id,dt)
```

2. For each id, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the a1c value set to missing. A two year gap would require two new rows, and so forth.

```
for (i in unique(x$id)){  
  for (j in 2:length(x$dt[x$id==i])){  
    dtnow = x$dt[x$id==i][j]  
    dtbefore = x$dt[x$id==i][j-1]  
    ytdiff = as.numeric(dtnow-dtbefore)/%/%365  
    while (ytdiff>0){  
      ytdiff = ytdiff-1  
      newrow = data.frame(id=i,dt=dtbefore+years(1),a1c=NA)  
      x = rbind(x,newrow)  
    }  
  }  
}
```

```
}
x = arrange(x,id,dt)
```

3. Create a new column `visit`. For each `id`, add the visit number. This should be 1 to `n` where `n` is the number of observations for an individual. This should include the observations created with missing `a1c` values.

```
x = x %>%
  group_by(id) %>%
  mutate(visit = 1:length(id))
```

4. For each `id`, replace missing values with the mean `a1c` value for that individual.

```
x = x %>%
  group_by(id) %>%
  mutate(a1c = replace_na(a1c,mean(a1c,na.rm = T)))
```

5. Print mean `a1c` for each `id`.

```
x %>%
  group_by(id) %>%
  summarize(mean_a1c=mean(a1c))
```

```
## # A tibble: 50 x 2
##       id mean_a1c
##   <int>   <dbl>
## 1     1     6.65
## 2     2     9.79
## 3     3     6.95
## 4     4     8.19
## 5     5     9.43
## 6     6     7.13
## 7     7     7.88
## 8     8     6.24
## 9     9     4.42
## 10    10     6.03
## # ... with 40 more rows
```

6. Print total number of visits for each `id`.

```
x %>%
  group_by(id) %>%
  summarize(visits=n())
```

```
## # A tibble: 50 x 2
##       id visits
##   <int>   <int>
## 1     1       7
## 2     2      16
## 3     3      13
## 4     4       9
## 5     5      14
## 6     6      11
## 7     7       7
## 8     8      12
## 9     9      15
## 10    10       8
```

```
## # ... with 40 more rows
```

7. Print the observations for id = 15.

```
x[x$id==15,]
```

```
## # A tibble: 10 x 4
## # Groups:   id [1]
##       id dt                a1c visit
##   <int> <dtm>              <dbl> <int>
## 1    15 2000-10-21 01:08:17  7.40     1
## 2    15 2001-08-08 14:23:08  5.90     2
## 3    15 2001-08-15 07:03:29  7.46     3
## 4    15 2002-03-15 21:23:10  5.33     4
## 5    15 2002-04-14 09:08:25  6.48     5
## 6    15 2002-10-10 18:27:43  8.14     6
## 7    15 2003-02-19 12:58:53  6.45     7
## 8    15 2003-03-02 06:58:10  7.43     8
## 9    15 2003-06-30 07:20:49  7.11     9
## 10   15 2004-01-22 20:30:42  5.67    10
```

Question 2

16 points

Install the `lexicon` package. Load the `sw_fry_1000` vector, which contains 1,000 common words.

```
data('sw_fry_1000', package = 'lexicon')
head(sw_fry_1000)
```

```
## [1] "the" "of" "to" "and" "a" "in"
```

1. Remove all non-alphabetical characters and make all characters lowercase. Save the result as `a`.

```
a = tolower(gsub("[^A-Za-z ]", "", sw_fry_1000))
```

Use vector `a` for the following questions. (2 points each)

2. How many words contain the string “ar”? 64

```
length(grep("ar",a))
```

```
## [1] 64
```

3. Find a six-letter word that starts with “l” and ends with “r”.

```
for (i in unique(a)){
  if (nchar(i)==6 &startsWith(i,"l")&endsWith(i,"r"))res = i}
res
```

```
## [1] "letter"
```

4. Return all words that start with “col” or end with “eck”.

```
res = NULL
for (i in unique(a)){
  if (startsWith(i,"col")|endsWith(i,"eck"))res = append(res,i)
}
res
```

```
## [1] "color" "cold" "check" "collect" "colony" "column" "neck"
```

5. Find the number of words that contain 4 or more adjacent consonants. Assume “y” is always a consonant. There are 8 words

```
str_subset(a, "[^aeiou]{4}")
```

```
## [1] "country"      "system"      "syllable"    "length"      "instrument"
## [6] "industry"     "symbol"      "supply"
```

6. Return all words with a “q” that isn’t followed by a “ui”.

```
str_subset(a, "q[u][^i]") # q must be followed by a u
```

```
## [1] "question" "equate"   "square"   "equal"     "quart"     "quotient"
```

7. Find all words that contain a “k” followed by another letter. Run the `table` command on the first character following the first “k” of each word.

```
a1 = str_subset(a, "k[^k]")
a2 = str_split(a1, "k")
res = NULL
for (i in 1:length(a2)){
  res = c(res, a2[[i]][2])
}
table(substr(res, 1, 1))
```

```
##
##  e  i  n  y
## 10  5  2  1
```

8. Remove all vowels. How many character strings are found exactly once? There are 581 strings found exactly once.

```
a3 = gsub("[aeiou]", "", a)
t = table(a3)
count = 0
for (i in 1:length(t)){
  if (as.numeric(t[i]) == 1) count = count + 1
}
count
```

```
## [1] 581
```

Question 3

3 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula.

```
haart_df = read.csv('haart.csv')[,c('death', 'weight', 'hemoglobin', 'cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here’s a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is “catching” the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(predvars, data, env): object 'death' not found>
```

What do you think is going on? Consider using `debug` to trace the problem.

The input of the function `response` cannot be directly assigned as `death` because it's not a valid value.

5 bonus points

Create a working function.

```
myfun <- function(dat, response) {
  y = dat[,response]
  other = dat[,!(colnames(dat) %in% response)]
  coef(summary(glm(y~other[,1]+other[,2]+other[,3], family=binomial(logit))))
}
myfun(haart_df, "death")
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## other[, 1]   -0.046210552 0.022556001 -2.048703 0.0404911395
## other[, 2]   -0.350642786 0.105064078 -3.337418 0.0008456055
## other[, 3]    0.002092582 0.001811959  1.154872 0.2481427160
```